SPEC-DRIVEN DEVELOPMENT

SpecSangam

ஸ்பெக் சங்கம்

My lightweight, measured Spec-Driven Development workflow, built with Claude.

The agent builds one phase at a time. You decide when it moves forward.

Spec
what to build

→

Agent
builds one phase

→

Human
approves the gate

Sangam means a coming together.

Sudarshan Ramanujam · Swipe →

HOW IT WORKS

Five steps, one phase at a time

1 Clarify the requirements

Claude asks 15 multiple-choice questions, one at a time, each with a recommended option and its trade-offs.

2 Approve the architecture

Scope, API contracts, constraints and exclusions go into `Architecture.md`.

3 Plan the work

`plan.md` defines phases and acceptance criteria. `tasks.md` has one section per phase, each task with a done-check.

4 Let CLAUDE.md drive the session

Each session reads only `CLAUDE.md` + `progress.md`, prints a 5-line summary and waits for **GO**. Then it opens just that phase.

5 Build, verify, hand off

Claude checks the acceptance criteria, records the handoff in `progress.md` and stops until I send **APPROVE**.

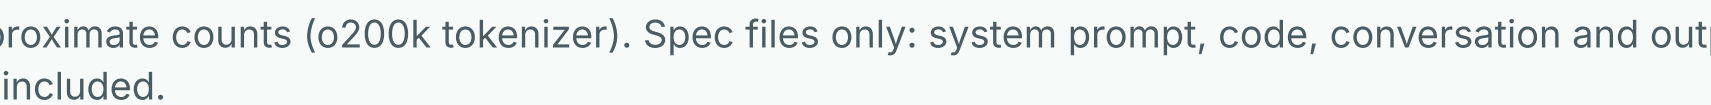
WHERE THE TOKEN COST DROPS

Load the phase, not the whole project

Spec tokens loaded at the start of a phase

Measured on my real Clausario spec files, at the start of Phase 5

Read all spec files **≈32,900 tokens**



CLAUDE.md + Architecture.md + plan.md + all of tasks.md + progress.md

SpecSangam phase start **≈11,100 tokens**

CLAUDE.md + progress.md + the Phase 5 section of tasks.md

~66% fewer spec tokens at the Phase 5 start.
Every phase: 60–83% (next page).

Less to reload

A short handoff and one phase's tasks, not the whole chat.

Less extra work

YAGNI keeps unapproved features out of the build.

Less rework

Each gate catches a wrong assumption early.

Approximate counts (o200k tokenizer). Spec files only: system prompt, code, conversation and output not included.

EVERY PHASE, MEASURED

Spec tokens loaded at each phase start

Phase being started	Read all spec files without SpecSangam	SpecSangam phase start with SpecSangam	This phase's tasks only part of SpecSangam total	Fewer spec tokens
P0	26,516	4,565	838	**83%**
P1	26,964	6,023	1,848	**78%**
P2	28,548	6,567	808	**77%**
P3	30,413	8,249	625	**73%**
P4	31,430	9,433	792	**70%**
P5	**32,916**	**11,099**	**972**	**66%**
P6	34,432	12,167	524	**65%**
P7	35,740	13,554	603	**62%**
P8	37,274	14,899	414	**60%**

All numbers are tokens (approximate). The saving falls as progress.md grows each phase. File-by-file example on the next page.

How it was measured

1. Real Clausario spec files.
2. progress.md counted up to the entries written before that phase.
3. Read all = CLAUDE.md + Architecture.md + plan.md + all of tasks.md + progress.md.
4. SpecSangam = CLAUDE.md + progress.md + that phase's section of tasks.md.
5. Saving = 1 − SpecSangam ÷ read all.

Limits

- o200k tokenizer; Claude's counts differ.
- Spec files only, not the whole session.
- Baseline is "read every spec file", not "no workflow".
- Files the ritual asks for; check with `/context`.

ONE EXAMPLE, FILE BY FILE

The Phase 5 start, with and without SpecSangam

File	Without SpecSangam	With SpecSangam
CLAUDE.md	3,000	3,000
progress.md notes from P0–P4	7,127	7,127
Architecture.md	7,995	*not read*
plan.md	7,228	*not read*
tasks.md	7,566 all 9 phases	972 Phase 5 section only
Total	**32,916**	**11,099**

Saving = 1 − (11,099 ÷ 32,916) = 1 − 0.34

≈ 66% fewer spec tokens

Architecture.md and plan.md stay closed, and tasks.md shrinks to one section. CLAUDE.md and progress.md are read either way.

WHAT KEEPS IT FOCUSED

Four practices, and what comes next

YAGNI

Build only what the phase names. No database until human review needed runs that pause and resume.

Approval gates

APPROVE after I review a phase, **GO** after the 5-line summary.

Caveman progress.md

Done, decided, open, next. After /clear, Claude resumes from these notes. My own terse format, not the plugin.

Docs every phase

README, Installation and Postman are updated before each gate.

NEXT: THE GOLDFISH STEP

A fresh agent reviews the plan before I approve it

Inspired by Google's Elephant–Goldfish Model. The reviewer reads only the files, with no conversation history, so it flags what the docs leave unsaid.

Elephant drafts · Goldfish reviews · Human approves

How do you keep an AI agent from building on a wrong assumption?

Tell me in the comments. Save this for your next build.

Builds on ideas popularised by Spec Kit, Kiro and GSD.

Sudarshan Ramanujam · SpecSangam